

STA 380 Assignment Part 2

Rohitashwa Chakraborty, Meha Mehta, Dipali Pandey, Sahitya Sundar Raj Vijayanagar

8/16/2021

Gtithub Link

Visual story telling part 1: green buildings

No we do not agree with the inference of the stat-guru.

While the guru's explanation sounds logical, the guru relies heavily on the assumption that the *green rating* of a building is the sole driver behind the rent of a property.

A simple linear regression run on the rent premium against green building indicates that it is statistically very significant.

```
##           Estimate Std. Error t value    Pr(>|t|)
## (Intercept)  0.7119434  0.1156967  6.153532 7.947381e-10
## green_rating1 2.4124581  0.3927573  6.142363 8.524150e-10
```

Coefficient of Green_Rating: 2.4125; P-Value: 8e-10

However, our model suffers from an extremely high bias and has a very low R^2 . This is a clear indicator that our model is inadequate.

Thus we try looking for confounding variables.

```
## 1 ) MLR, green-building rating with CS_PropertyID : Non Confounding
## 2 ) MLR, green-building rating with cluster : Non Confounding
## 3 ) MLR, green-building rating with size : Non Confounding
## 4 ) MLR, green-building rating with empl_gr : Non Confounding
## 5 ) MLR, green-building rating with leasing_rate : Non Confounding
## 6 ) MLR, green-building rating with stories : Non Confounding
## 7 ) MLR, green-building rating with age : Non Confounding
## 8 ) MLR, green-building rating with renovated : Non Confounding
## 9 ) MLR, green-building rating with class_a : Confounding
## 10 ) MLR, green-building rating with class_b : Non Confounding
## 11 ) MLR, green-building rating with net : Non Confounding
## 12 ) MLR, green-building rating with amenities : Non Confounding
## 13 ) MLR, green-building rating with cd_total_07 : Non Confounding
## 14 ) MLR, green-building rating with hd_total07 : Non Confounding
## 15 ) MLR, green-building rating with total_dd_07 : Non Confounding
## 16 ) MLR, green-building rating with Precipitation : Non Confounding
## 17 ) MLR, green-building rating with Gas_Costs : Non Confounding
## 18 ) MLR, green-building rating with Electricity_Costs : Non Confounding
```


Therefore, It seems that whether a building is rated as a *Class-A* listing is actually an underlying confounding variable with the *Rent* and *Green Ratings*

THis is not surprising since a *Class-A* listing is the most desirable property and will be superior to its neighbourhood competition in terms of not only amenities and services but will also be technologically superior and therefore will have lower costs. All these factors will raise the price and also increase the chances that the property qualifies as a *Green* building.

```
##           Estimate Std. Error  t value
## (Intercept) -0.7401308  0.1405214 -5.267032
## green_rating1 0.6688835  0.3979776  1.680706
## class_a1     4.0091931  0.2286928 17.530914
```

The T-statistic shows that at a 95% confidence, *Green Rating* is in fact not statistically significant in determining the rent!

It is the Class-A rating which determines the rent instead!!




The medians (black lines on the box-plots) and means (the red stars) of green and non-green building revenue per square foot show that green buildings have higher revenue. The non-green buildings have a lower mean than the green buildings

Visualizations



##Inference

- There is a correlation between rent and the cluster rent
- The size of the rental space in the building is also correlated with the Rent
- A Class buildings appear to be younger
- Age does not seem to have a high correlation with rent
- Class A buildings have higher rent

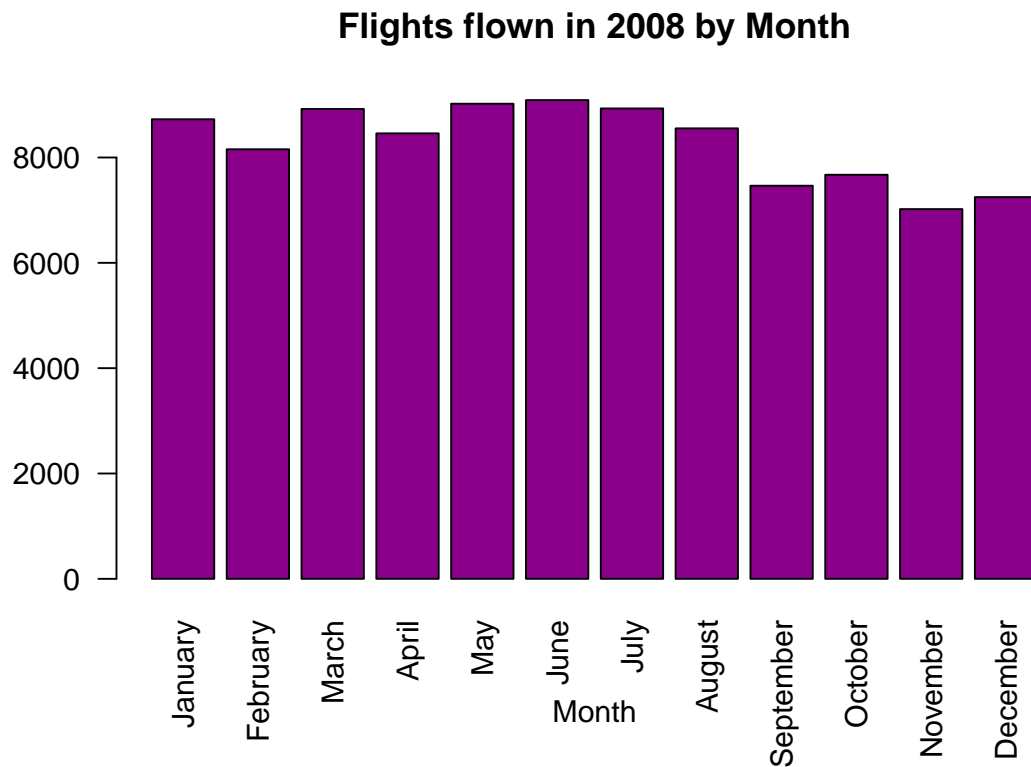




Visual story telling part 2: flights at ABIA

Flights flown by Month

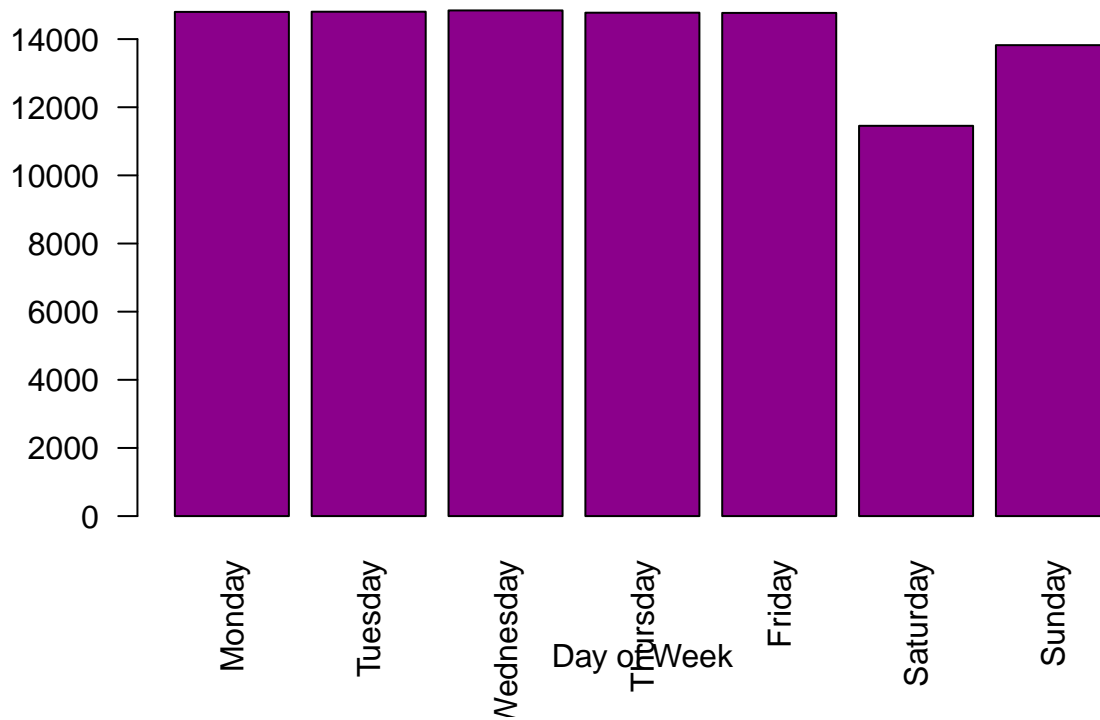
- Number of flights flying during the last 4 months of the year much lesser than other months



Flights flown by Day of week

- Comparatively lesser flights flown on Saturdays than other days


Flights flown in 2008 by Day of Week



Flights flown by Carrier

- Southwest Airlines leads in the number of flights flown in 2008, followed by American Airlines


Flights flown in 2008 by Unique Carrier

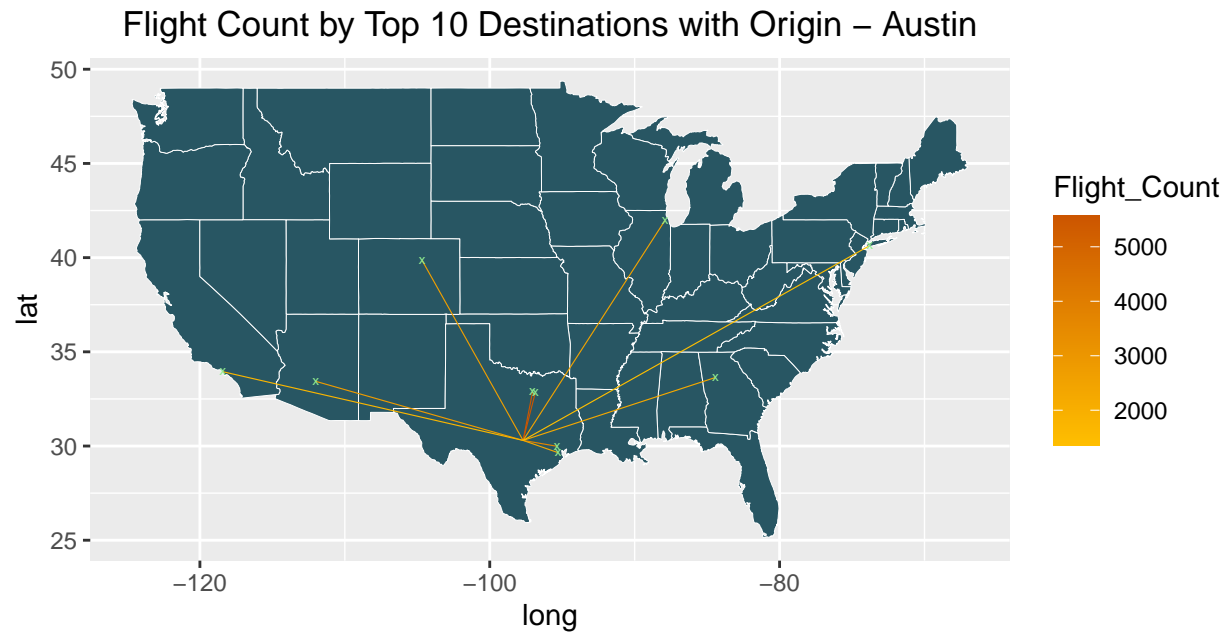


Top 20 destinations in 2008 for flights flown from Austin

- Top destinations: Dallas, Houston, Phoenix, Denver

Flights flown from Austin in 2008 by Top 20 Destination







Top 20 Origins in 2008 for flights flown to Austin

- Top Origins: Dallas, Houston, Phoenix, Denver


Flights flown to Austin in 2008 by Top 20 Origins



Plotting average arrival delays from different origins




Arrival Delays Chart



- Flight arrival delay statistics

Departure Delays Chart

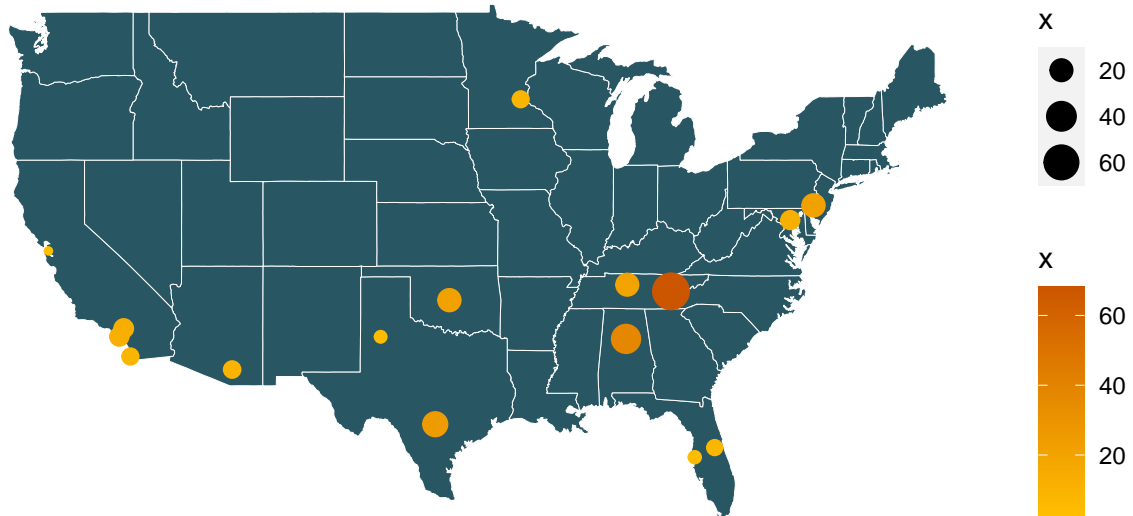


- Flight departure delay statistics

Flight Delays by Origin Airport

- To Determine which airports have maximum average delays for flights to ABIA

Delays by Origin Airport



Portfolio modeling

Three ETF portfolios of different asset classes were invested into to test their respective four-trading-week Value at Risk at the 5% level using past five years of data (starting from 2016-01-01). The three classes picked were HY Bonds, Corporate (Investment Grade) Bonds and Government Bonds. \$100,000 were allocated for each of the three portfolios containing five ETFs with equal weights (20% in each ETF) which would rebalance daily i.e. each of the ETFs in a portfolio will always comprise 20% of the total wealth at the beginning of each trading day.

High Yield Bonds Portfolio

- For this portfolio, the US traded high yield bond ETFs were selected. This portfolio contains the follow ETFs:
- HYG: Beta = 0.37; Weight = 20%
- JNK: Beta = 0.40; Weight = 20%
- USHY: Beta = 0.40; Weight = 20%
- SRLN: Beta = 0.34; Weight = 20%
- BKLN: Beta = 0.37; Weight = 20%
- Weighted average Beta of the portfolio is 0.376.

- This portfolio was selected as High Yield Bonds ETFs offer investors exposure to debt issued by below investment grade corporations. These ETFs invest in junk bonds, senior loans, as well as international below investment grade debt. These products have offered numerous ways for investors to take advantage of this space. High yields can be a great addition to a yield-hungry portfolio, as they can offer yields into the double digits for those willing to take on the risks that come along with it. The high returns come from riskier bond choices who have to pay out higher ratios to compensate investors for high risks. This means that the holdings of these ETFs will have higher chances of defaults, and could potentially leave investors out to dry. But for those who have done their homework on the holdings of a particular “junk” bond fund have the ability to generate strong returns from these powerful products.


Investment Grade Bonds Portfolio

- For this portfolio, the US traded investment grade bond ETFs were selected.
- VCIT: Beta = 0.18; Weight = 20%
- LQD: Beta = 0.19; Weight = 20%
- VCSH: Beta = 0.08; Weight = 20%
- IGSB: Beta = 0.08; Weight = 20%
- IGIB: Beta = 0.20; Weight = 20%
- Weighted average Beta of the portfolio is 0.146.
- This portfolio was selected to gain exposure to investment grade corporate bonds, which provides low risk, stable return profile as they are invested in investment-grade credit (highly) rated corporates. However, the returns are greater than the government bonds. Bonds included in these funds can feature varying maturities and are issued by companies from multiple industries.

Government Bonds Portfolio

- For this portfolio, the US traded government bond ETFs were selected. This portfolio contains the follow ETFs:
- SHY: Beta = -0.03; Weight = 20%
- TLT: Beta = -0.25; Weight = 20%
- GOVT: Beta = -0.09; Weight = 20%
- SHV: Beta = -0.01; Weight = 20%
- IEF: Beta = -0.13; Weight = 20%
- Weighted average Beta of the portfolio is -0.102
- This portfolio was selected to gain exposure to fixed income securities issued by government agencies. Bonds featured in these ETFs include U.S. Treasuries of varying maturities, floating rate Treasury bonds, and TIPS. The aim for this portfolio is to earn stable return while taking a limited amount of risk. The bonds are mostly either backed by the US Treasury or state governments, the default risk is near zero implying an extremely low risk as indicated by their negative Betas.

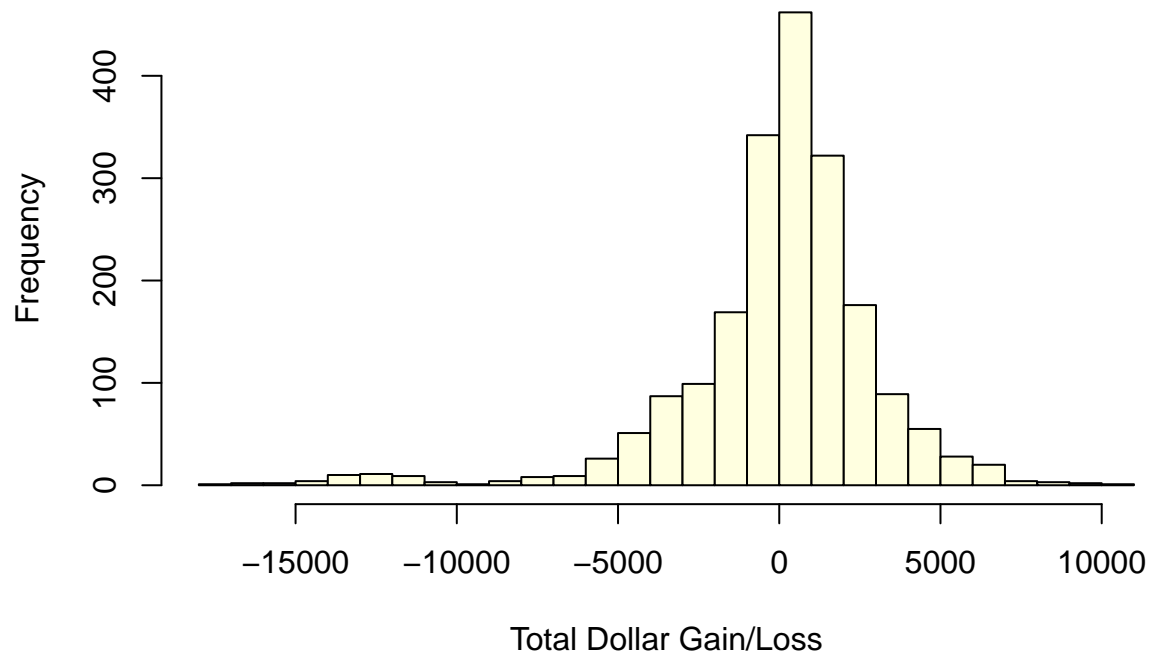
High Yield Bonds Portfolio_Wealth



```
## [1] 100009.3
```

```
## [1] 9.349063
```

High Yield Bonds Portfolio_Gain/Loss

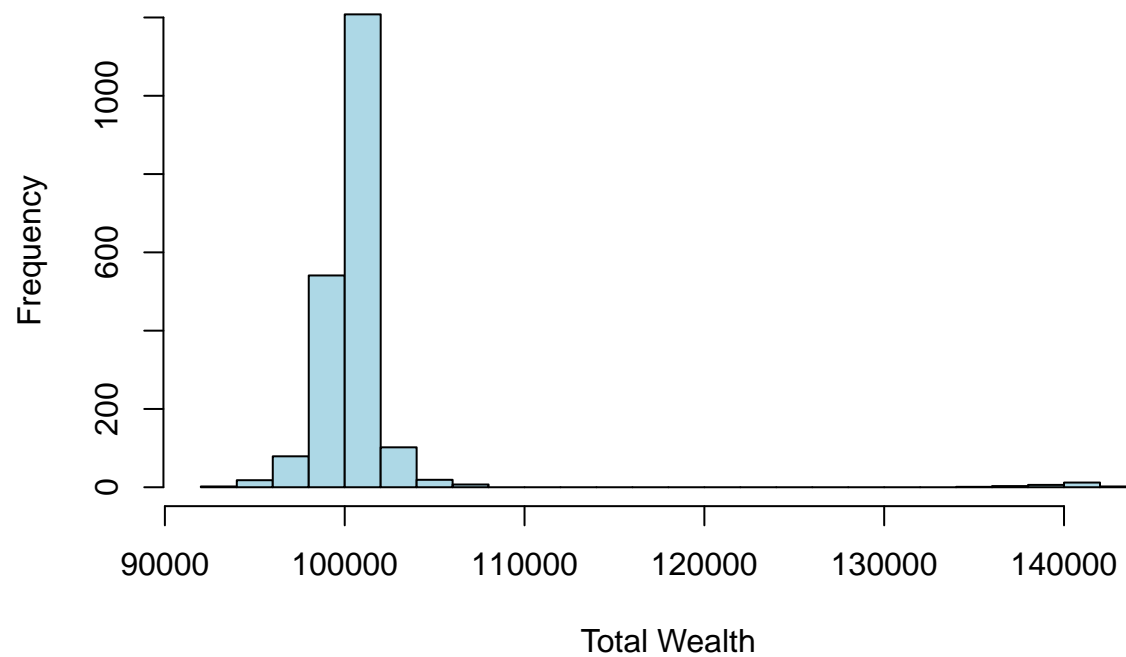


5%
-4688.555

5%
-0.04688555

VaR=4.688555%

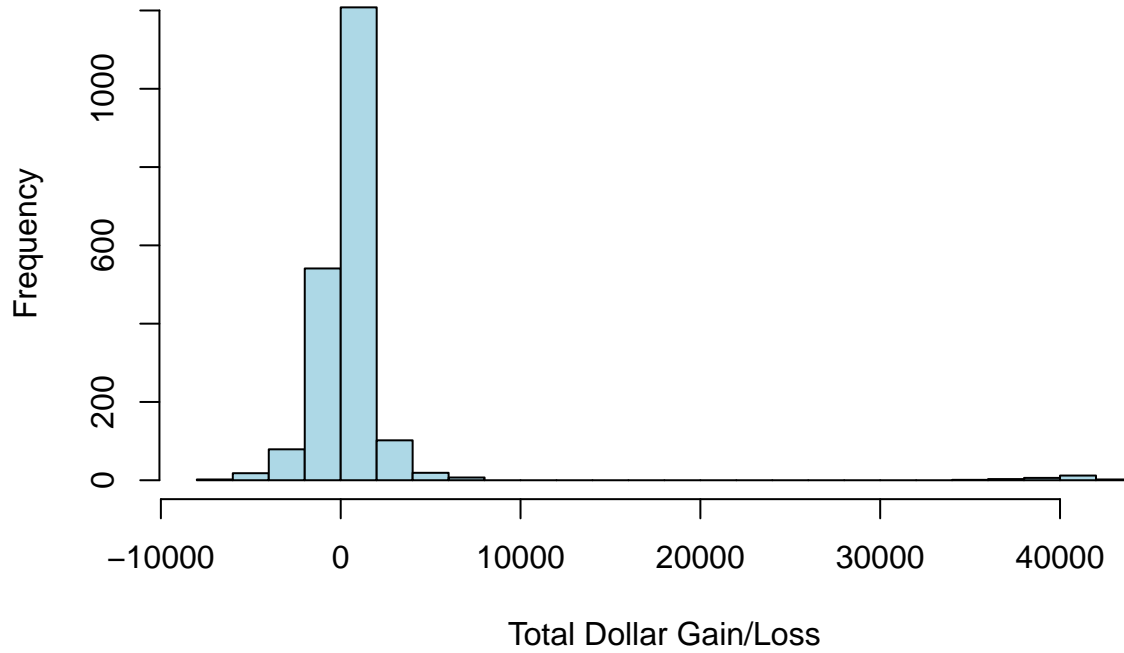
Investment Grade Bonds Portfolio_Wealth



```
## [1] 100836.5
```

```
## [1] 836.5093
```

Investment Grade Bonds Portfolio_Gain/Loss




5%
-1956.064

5%
-0.01956064

VaR=-1.956064%

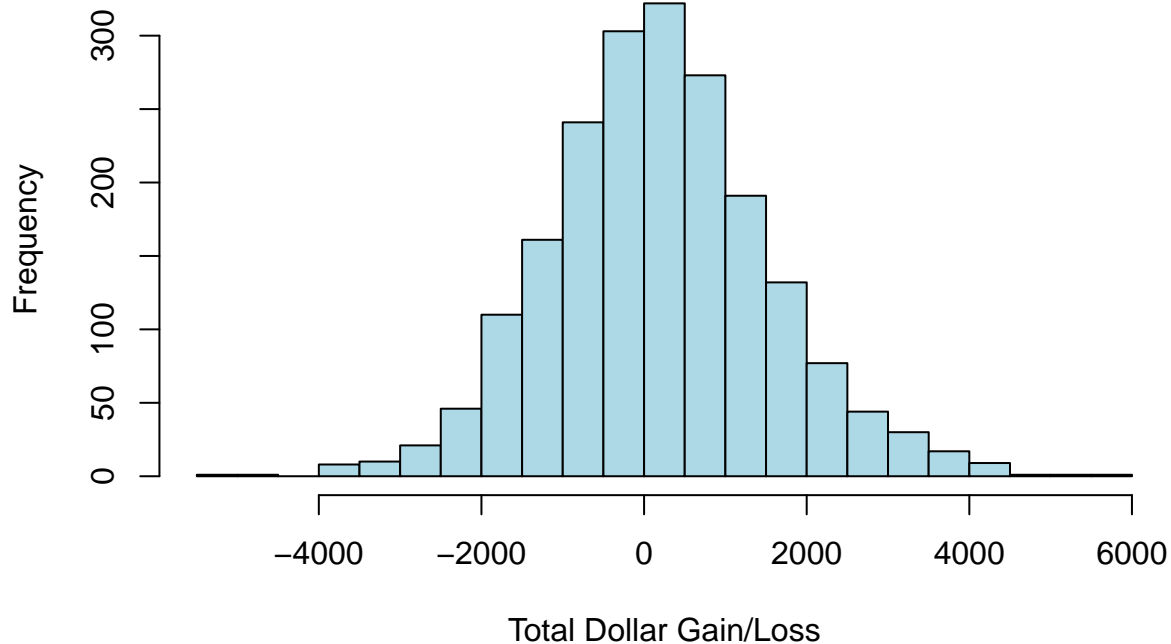
Government Bonds Portfolio_Wealth



```
## [1] 100188.4
```

```
## [1] 188.3696
```

Government Bonds Portfolio_Gain/Loss



```
##          5%
## -1948.256
```

```
##          5%
## -0.01948256
```

```
-0.01948256 -1.948256%
```

- 5% Value at Risk for the three portfolios were:
- High Yield Bond Portfolio: -\$6,997.02 or -4.69%
- Investment Grade Bond Portfolio: -1.96%
- Government Bond Portfolio: -1.95%
- Investing \$100,000 in each of the portfolios for 20 trading days with daily rebalancing, there is a 5% chance for our total gain/loss for (1) the High Yield Bond Portfolio to be less than -4.69%, (2) the Investment Grade Bond Portfolio to be less than -1.96%, (3) the Government Bond Portfolio to be 1.95%. It is clearly seen that the High Yield portfolio has chances of higher loss due to higher risk involved. The Investment Grade and Government Bond Portfolios appear to have comparatively low and similar loss chances as they are low risk, with the Government one having the least Value at Risk.

Market segmentation


Market Segmentation

Data Pre-Processing

In our Dataset, we observe the following:

- X column is essentially the Turk's ID. Thus, this column has been set as the row index and dropped consequently
- Mean-normalised all columns and scaled them to a 0-1 range (*Performed Z-transformation*) to make the measurements comparable

Elbow method




The Elbow Plot above indicates a kink at $K = 14$. This could be a good estimate of the number of clusters we go choose.

However, This method is too *visual* and the differences are not clearly discernible. Thus, let us re-validate our results with the calculation of the **GAP** metric.

Gap Index

clusGap(x = NutrientH20, FUNcluster = kmeans, K.max = 20, B = 3, nstart = 5, iter.max = 1000)





```
##          logW    E.logW      gap      SE.sim
## [1,]  9.666055 10.99032  1.324268 0.0008224680
## [2,]  9.611718 10.95224  1.340521 0.0008446234
## [3,]  9.583689 10.93412  1.350427 0.0004244282
## [4,]  9.561569 10.92045  1.358884 0.0007422304
## [5,]  9.538090 10.91074  1.372646 0.0010602093
## [6,]  9.518623 10.90187  1.383246 0.0008250187
## [7,]  9.498126 10.89502  1.396896 0.0003843231
## [8,]  9.482833 10.88803  1.405194 0.0004427242
## [9,]  9.469945 10.88255  1.412604 0.0004554329
## [10,] 9.454182 10.87747  1.423289 0.0007160178
## [11,] 9.447812 10.87315  1.425342 0.0008341239
## [12,] 9.433375 10.86909  1.435713 0.0005789388
## [13,] 9.429150 10.86523  1.436080 0.0007802581
## [14,] 9.414951 10.86167  1.446714 0.0006919865
## [15,] 9.411458 10.85793  1.446474 0.0006214817
## [16,] 9.402572 10.85456  1.451986 0.0003996070
## [17,] 9.395897 10.85187  1.455969 0.0005255290
## [18,] 9.389830 10.84893  1.459097 0.0003837243
## [19,] 9.384525 10.84584  1.461313 0.0005703963
## [20,] 9.381393 10.84354  1.462150 0.0004201965
```

##

```
##  
## Optimum cluster size is: 14
```

From the above plots and output statistic, we see that the local optima of the Cluster Gap plot is achieved at $K = 14$ clusters. Thus, we go ahead with dividing the consumer base in 14 clusters for targeted marketing.

Running K-Means with 14 clusters




```
##
## Cluster Whithinness: 11122.7847934064 10237.0217657893 18304.1104205402 14996.9611695742 6190.420127

##
## Total Whithinness: 170419.12158895

##
## Cluster Betweenness: 113296.87841105
```

Author attribution

- Building a dictionary from training data (half of ReutersC50)

```
## <<DocumentTermMatrix (documents: 2500, terms: 32669)>>
## Non-/sparse entries: 619802/81052698
## Sparsity           : 99%
## Maximal term length: 40
## Weighting           : term frequency (tf)
```

```
## <<DocumentTermMatrix (documents: 2500, terms: 32570)>>
## Non-/sparse entries: 537861/80887139
## Sparsity           : 99%
## Maximal term length: 40
## Weighting           : term frequency (tf)
```

```
## <<DocumentTermMatrix (documents: 2500, terms: 3393)>>
## Non-/sparse entries: 422971/8059529
## Sparsity           : 95%
## Maximal term length: 40
## Weighting           : term frequency (tf)
```

Pre-processing steps:

- Convert to lowercase
- Numbers removed
- Punctuation removed
- Unnecessary white space trimmed

Data reduced to 2500 documents with 32,669 terms

- Stop words removed

Data reduced to 32,570 terms.

- Sparse terms removed

Data reduced to 3393 terms.

- Word counts changed to TF-IDF weights.
- Process repeated for testing data. 3448 terms in the testing data compared to 3393 terms in the training data. We must consider how to treat these new words.

```
## <<DocumentTermMatrix (documents: 2500, terms: 33472)>>
## Non-/sparse entries: 628611/83051389
## Sparsity           : 99%
## Maximal term length: 45
## Weighting          : term frequency (tf)
```

```
## <<DocumentTermMatrix (documents: 2500, terms: 33373)>>
## Non-/sparse entries: 545286/82887214
## Sparsity           : 99%
## Maximal term length: 45
## Weighting          : term frequency (tf)
```

```
## <<DocumentTermMatrix (documents: 2500, terms: 3448)>>
## Non-/sparse entries: 428509/8191491
## Sparsity           : 95%
## Maximal term length: 39
## Weighting          : term frequency (tf)
```


- 55 new terms observed in the testing data
- new terms ignored since they constitute less than 2% of the training data

```
## <<DocumentTermMatrix (documents: 2500, terms: 3393)>>
## Non-/sparse entries: 379314/8103186
## Sparsity           : 96%
## Maximal term length: 40
## Weighting          : term frequency - inverse document frequency (normalized) (tf-idf)
```

The training set now has 3393 predictors. Principal Component Analysis (PCA) is performed to reduce the number of predictors as follows:

- Eliminate columns that have 0 entries. This eliminates columns where the term is not found in any data in the test or train set. This ensures the term lists are identical and only use columns common between train and test data, leaving 8,317,500 elements in both the train and test matrices.
- Once the data is in the same format, use PCA to choose the number of principal components.

Summary of Principal Component Variance Analysis



- Choice of number of components is based on desired level of goodness of fit
- At 1000 principal components ~80% of the variance is explained
- Data pre-processing is complete and models can be run

Four models run to predict author attribution:

- KNN
- Random Forest
- Naive Bayes
- Logistic regression

KNN

KNN model run with $k = 1$, other values of k did not improve the testing or training accuracy.

```
## [1] 0.326
```

- Maximum testing accuracy obtained with KNN : 32.6%

Random Forest

- Random Forest model run with $m=31$ (square root of 1000, the number of total variables).

```
## [1] 0.784
```

- Maximum testing accuracy obtained with Random Forest : 78.4%
- Large improvement over KNN

Naive Bayes

- Naive Bayes model used to predict the testing data from a training model.

```
## [1] 0.3144
```

- Maximum testing accuracy obtained with Naive Bayes : 31.44%
- Reduction in accuracy from KNN

Multinomial Logistic Regression

- Lastly, multinomial logistic regression performed with only the first 18 Principal Components

```
## # weights: 1000 (931 variable)
## initial value 9780.057514
## iter 10 value 3978.610720
## iter 20 value 3027.473914
## iter 30 value 2744.784042
## iter 40 value 2654.854912
## iter 50 value 2552.955018
## iter 60 value 2433.970187
## iter 70 value 2345.527733
## iter 80 value 2305.565551
## iter 90 value 2285.592449
## iter 100 value 2269.532712
## final value 2269.532712
## stopped after 100 iterations
```


```
## [1] 0.4652
```

- Maximum testing accuracy obtained with Logistic Regression : 46.52%
- Better than KNN or Naive Bayes
- Lower than Random Forest

Comparing the accuracy of the various models:

```
##           models accuracy
## 1      Random Forest    78.40
## 2 Logistic Regression    46.52
## 3              KNN      32.60
## 4      Naive Bayes      31.40
```

Summary of Attribution Models Classification Accuracy



In summary, the random forest model has the highest classification accuracy of ~78.4% on the testing dataset.

Association rule mining


Load the dataset from 'groceries.txt' as transactions of format basket. As shown, there are 2159 transactions with just one item in the basket, 1643 transactions with two items in the basket, and so on.

```
## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 169 columns (items) and a density of 0.02609146
##
## most frequent items:
##      whole milk other vegetables      rolls/buns      soda
##           2513           1903           1809           1715
##           yogurt      (Other)
##           1372           34055
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16
## 2159 1643 1299 1005  855  645  545  438  350  246  182  117  78  77  55  46
##      17     18     19     20     21     22     23     24     26     27     28     29     32
##      29     14     14      9     11      4      6      1      1      1      1      3      1
##
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    1.000   2.000   3.000   4.409   6.000   32.000
##
## includes extended item information - examples:
##           labels
## 1 abrasive cleaner
## 2 artif. sweetener
## 3  baby cosmetics
```

Item Frequency Plot

Create the Item frequency plot to indicate the frequencies for different bought items. As shown, whole milk is the most frequently bought grocery item, followed by other vegetables and rolls/buns.



Association mining rules

Creating association mining rules by randomly selecting minimum support as 0.001, confidence as 0.8, and max length=10, results in 410 rules.

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##           0.8    0.1    1 none FALSE                TRUE         5   0.001    1
## maxlen target  ext
```

```

##      10  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 9
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [157 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 6 done [0.01s].
## writing ... [410 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

```

	lhs	rhs	support	confidence	coverage	lift	count
## [1]	{liquor, red/blush wine}	=> {bottled beer}	0.001931876	0.9047619	0.002135231	11.235269	
## [2]	{cereals, curd}	=> {whole milk}	0.001016777	0.9090909	0.001118454	3.557863	
## [3]	{cereals, yogurt}	=> {whole milk}	0.001728521	0.8095238	0.002135231	3.168192	
## [4]	{butter, jam}	=> {whole milk}	0.001016777	0.8333333	0.001220132	3.261374	
## [5]	{bottled beer, soups}	=> {whole milk}	0.001118454	0.9166667	0.001220132	3.587512	
## [6]	{house keeping products, napkins}	=> {whole milk}	0.001321810	0.8125000	0.001626843	3.179840	
## [7]	{house keeping products, whipped/sour cream}	=> {whole milk}	0.001220132	0.9230769	0.001321810	3.612599	
## [8]	{pastry, sweet spreads}	=> {whole milk}	0.001016777	0.9090909	0.001118454	3.557863	
## [9]	{curd, turkey}	=> {other vegetables}	0.001220132	0.8000000	0.001525165	4.134524	
## [10]	{rice, sugar}	=> {whole milk}	0.001220132	1.0000000	0.001220132	3.913649	

Looking at the above result, 100% of customers who bought {rice,sugar} also bought whole milk. Similarly, 90.48% of customers who bought {liquor, red/blush wine} also bought bottled beer.

Trying more stricter rules with conf=0.9 and shorter rules with maxlen=3, results in only 10 rules as shown below:

```

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.9      0.1      1 none FALSE              TRUE          5      0.001      1
## maxlen target  ext
##      3 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose

```

```
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 9
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [157 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 done [0.00s].
## writing ... [10 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

##      lhs                                rhs                support confidence    coverage    lift coun
## [1] {liquor,                                => {bottled beer}    0.001931876  0.9047619 0.002135231 11.235269
##      red/blush wine}
## [2] {cereals,                                => {whole milk}    0.001016777  0.9090909 0.001118454  3.557863
##      curd}
## [3] {bottled beer,                            => {whole milk}    0.001118454  0.9166667 0.001220132  3.587512
##      soups}
## [4] {house keeping products,                  => {whole milk}    0.001220132  0.9230769 0.001321810  3.612599
##      whipped/sour cream}
## [5] {pastry,                                  => {whole milk}    0.001016777  0.9090909 0.001118454  3.557863
##      sweet spreads}
## [6] {rice,                                    => {whole milk}    0.001220132  1.0000000 0.001220132  3.913649
##      sugar}
## [7] {bottled water,                            => {whole milk}    0.001220132  0.9230769 0.001321810  3.612599
##      rice}
## [8] {canned fish,                             => {whole milk}    0.001118454  1.0000000 0.001118454  3.913649
##      hygiene articles}
## [9] {grapes,                                  => {other vegetables} 0.001118454  0.9166667 0.001220132  4.737476
##      onions}
## [10] {hard cheese,                             => {other vegetables} 0.001118454  0.9166667 0.001220132  4.737476
##      oil}
```

Visualizations

Considering the 410 association rules created by minimum support as 0.001, confidence as 0.8, and max length=10, the next step included removing subsets of larger rules, which resulted in a total of 319 rules.

```
## [1] 91
```


```
##      lhs                                rhs                support confidence    coverage    lift coun
## [1] {liquor,                                => {bottled beer}    0.001931876  0.9047619 0.002135231 11.235269
##      red/blush wine}
## [2] {cereals,                                => {whole milk}    0.001016777  0.9090909 0.001118454  3.557863
##      curd}
## [3] {cereals,                                => {whole milk}    0.001728521  0.8095238 0.002135231  3.168192
##      yogurt}
## [4] {butter,                                  => {whole milk}    0.001016777  0.8333333 0.001220132  3.261374
##      jam}
## [5] {bottled beer,                            => {whole milk}    0.001118454  0.9166667 0.001220132  3.587512
##      soups}
```




```

## [6] {house keeping products,
##      napkins}          => {whole milk}      0.001321810  0.8125000 0.001626843  3.179840
## [7] {house keeping products,
##      whipped/sour cream} => {whole milk}      0.001220132  0.9230769 0.001321810  3.612599
## [8] {pastry,
##      sweet spreads}     => {whole milk}      0.001016777  0.9090909 0.001118454  3.557863
## [9] {curd,
##      turkey}            => {other vegetables} 0.001220132  0.8000000 0.001525165  4.134524
## [10] {rice,
##       sugar}            => {whole milk}      0.001220132  1.0000000 0.001220132  3.913649

```





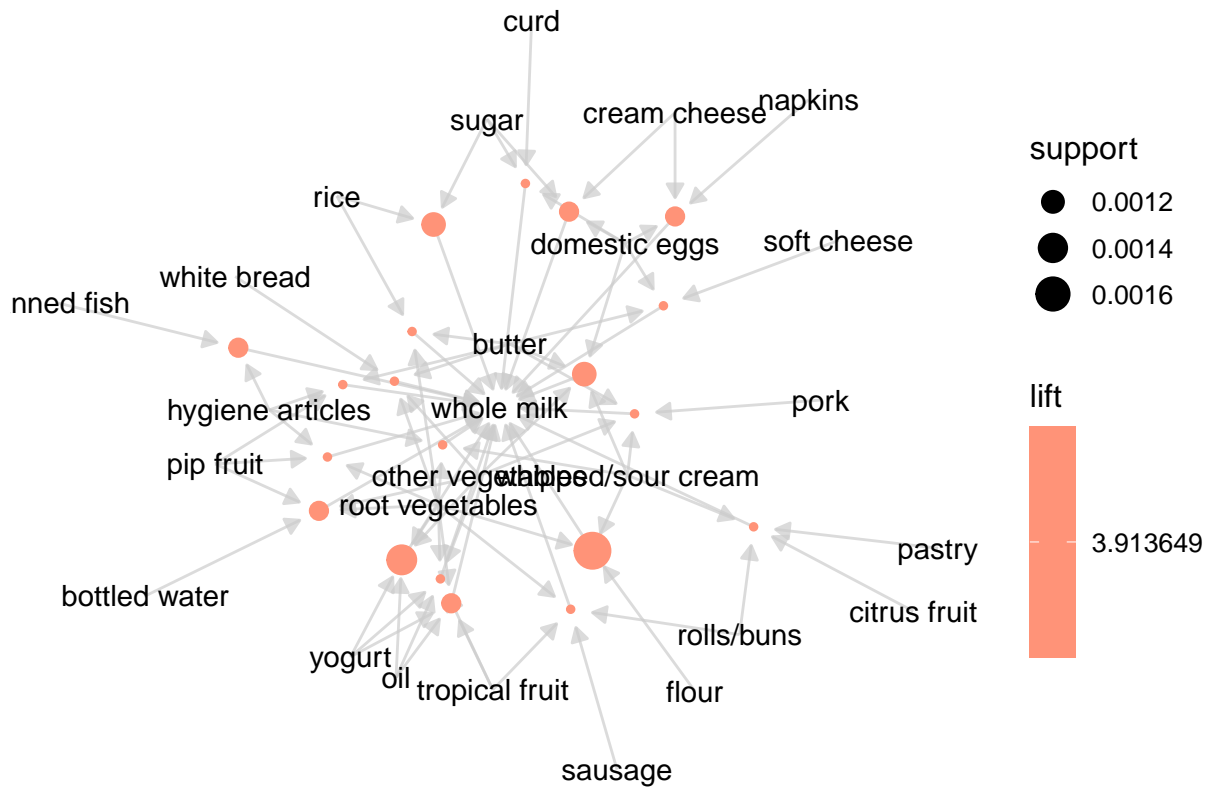
Recommendation rules

Looking at the most bought item in the item frequency list, i.e.m, 'whole milk', it is possible to find the items that are most likely to be bought before buying whole milk by using appearance. Further, using $\text{conf}=1$, will indicate the items where 100% of customers bought whole milk, after buying these items.

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          1    0.1    1 none FALSE                TRUE      5    0.001    1
## maxlen target  ext
##      10  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 9
##
## set item appearances ...[1 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [157 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 6 done [0.01s].
```

```
## writing ... [20 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

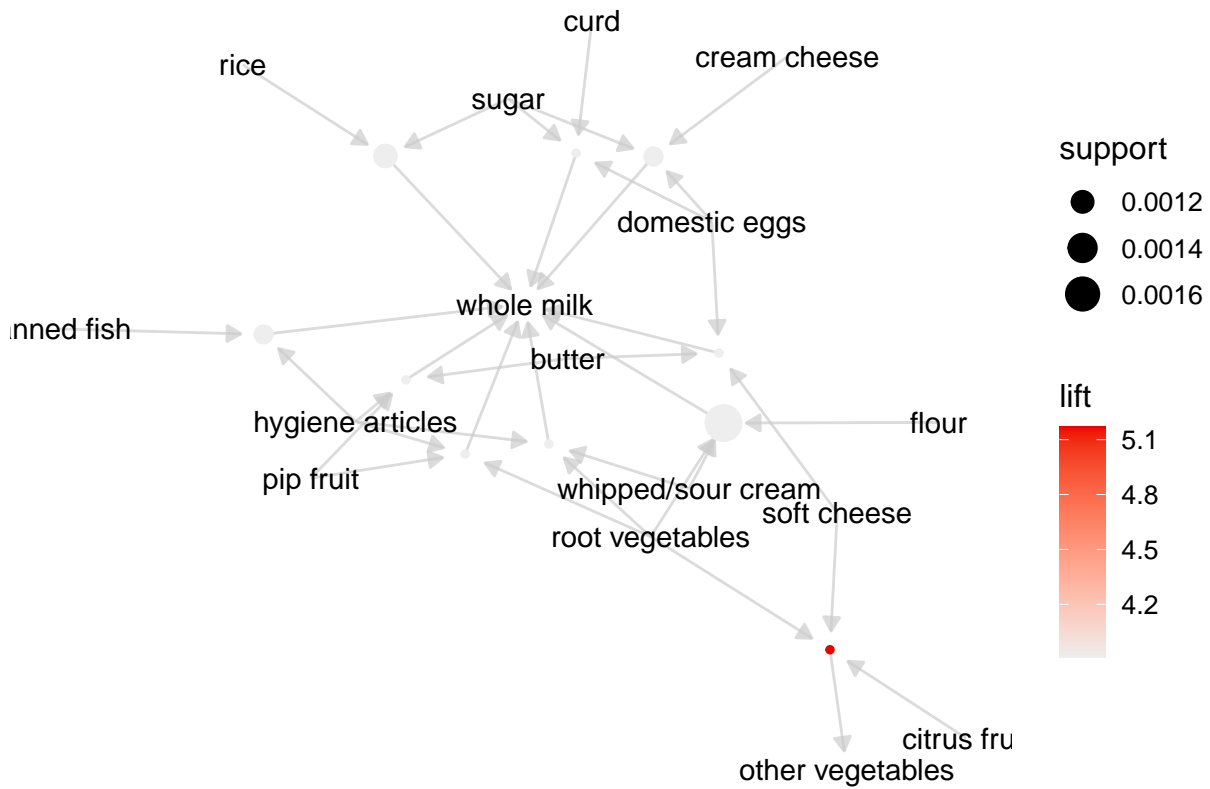
	lhs	rhs	support	confidence	coverage	lift	count
## [1]	{rice, sugar}	=> {whole milk}	0.001220132	1	0.001220132	3.913649	12
## [2]	{canned fish, hygiene articles}	=> {whole milk}	0.001118454	1	0.001118454	3.913649	11
## [3]	{butter, rice, root vegetables}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10
## [4]	{flour, root vegetables, whipped/sour cream}	=> {whole milk}	0.001728521	1	0.001728521	3.913649	17
## [5]	{butter, domestic eggs, soft cheese}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10
## [6]	{butter, hygiene articles, pip fruit}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10
## [7]	{hygiene articles, root vegetables, whipped/sour cream}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10
## [8]	{hygiene articles, pip fruit, root vegetables}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10
## [9]	{cream cheese, domestic eggs, sugar}	=> {whole milk}	0.001118454	1	0.001118454	3.913649	11
## [10]	{curd, domestic eggs, sugar}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10



As shown above, 100% of customers who bought {rice,sugar}, {canned fish,hygiene articles}, {butter,rice,root vegetables}, etc. have bought whole milk, Similarly, we find 20 such antecedents for whole milk.

Sorting rules based on confidence

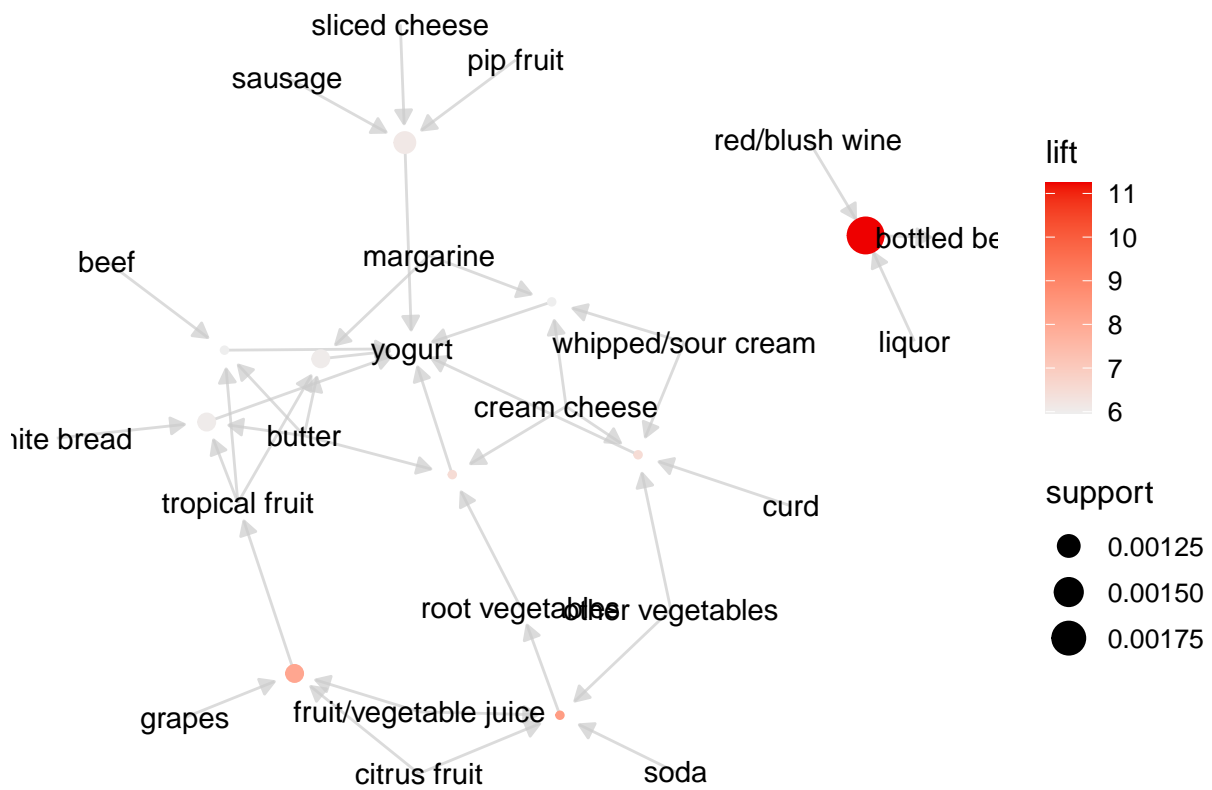
##	lhs	rhs	support	confidence	coverage	lift	count
## [1]	{rice, sugar}	=> {whole milk}	0.001220132	1	0.001220132	3.913649	12
## [2]	{canned fish, hygiene articles}	=> {whole milk}	0.001118454	1	0.001118454	3.913649	11
## [3]	{flour, root vegetables, whipped/sour cream}	=> {whole milk}	0.001728521	1	0.001728521	3.913649	17
## [4]	{butter, domestic eggs, soft cheese}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10
## [5]	{citrus fruit, root vegetables, soft cheese}	=> {other vegetables}	0.001016777	1	0.001016777	5.168156	10
## [6]	{butter, hygiene articles, pip fruit}	=> {whole milk}	0.001016777	1	0.001016777	3.913649	10



A confidence of 1 indicates that whenever the items on the antecedent are bought, 100% of customers bought the item(s) on the consequent.

Sorting rules based on lift

##	lhs	rhs	support	confidence	coverage	lift	count
## [1]	{liquor,	=> {bottled beer}	0.001931876	0.9047619	0.002135231	11.235269	19
##	red/blush wine}						
## [2]	{citrus fruit,	=> {root vegetables}	0.001016777	0.9090909	0.001118454	8.340400	10
##	fruit/vegetable juice,						
##	other vegetables,						
##	soda}						
## [3]	{citrus fruit,	=> {tropical fruit}	0.001118454	0.8461538	0.001321810	8.063879	11
##	fruit/vegetable juice,						
##	grapes}						
## [4]	{butter,	=> {yogurt}	0.001016777	0.9090909	0.001118454	6.516698	10
##	cream cheese,						
##	root vegetables}						
## [5]	{cream cheese,	=> {yogurt}	0.001016777	0.9090909	0.001118454	6.516698	10
##	curd,						
##	other vegetables,						
##	whipped/sour cream}						
## [6]	{pip fruit,	=> {yogurt}	0.001220132	0.8571429	0.001423488	6.144315	12
##	sausage,						
##	sliced cheese}						



A rule with lift of 11.23 for {liquore,red/blush wine}->{bottled beer} indicates that the items in the antecedent and consequent are ~11 times more likely to be bought together than bought individually.